



# **Project #1**

## **Multinomial logistic regression**

**Student :** David Albert

# Contents

<b>I</b>	<b>Task 1: Multinomial logistic regression</b>	<b>3</b>
1	Code explanation and results . . . . .	3
1.1	Code explanation . . . . .	3
1.2	Outputs . . . . .	4
2	Derivations of the optimal estimator and the optimal average MSE . . . . .	5
2.1	Integration of the 1st part . . . . .	5
2.2	Integration of the 2nd part . . . . .	5
2.3	Integration of the 3rd part . . . . .	6
2.4	Minimization of average MSE over all $\theta$ . . . . .	6
<b>II</b>	<b>Task 2 : Rain or shine</b>	<b>8</b>
1	Explanation and outputs of <i>project1_task2_generate.py</i> . . . . .	8
1.1	Explanation . . . . .	8
1.2	Outputs . . . . .	9
2	Explanation and outputs of <i>project1_task2.py</i> . . . . .	9
2.1	Explanations . . . . .	9
2.2	Outputs . . . . .	11

# Part I

## Task 1: Multinomial logistic regression

### 1 Code explanation and results

#### 1.1 Code explanation

Creating the computational graph using softmax function as output and so *softmax cross entropy with logits* as loss function.

```
x = tf.placeholder(tf.float32, shape=[None, 3])
y_true = tf.placeholder(tf.float32, shape=[None, 3])

w = tf.Variable(tf.truncated_normal([3, 3], stddev=0.5))
b = tf.Variable(tf.constant(0.1, shape=[3]))
z = tf.matmul(x, w) + b
y_pred = tf.nn.softmax(z)

loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(
        labels=y_true, logits=z))
train = tf.train.GradientDescentOptimizer(0.1).minimize(loss)
mse = tf.reduce_mean((y_pred - y_true)**2)
```

At each epoch, we create  $n = \text{mini\_batch\_size}$  random samples of  $\theta = \{\theta_0, \theta_1, \theta_2\}$  from the uniform distribution on the simplex  $\theta_0 + \theta_1 + \theta_2 = 1$ . It is similar to sampling from the Dirichlet(1,1,1) distribution.

```
# Sampling theta=(theta0,theta1,theta2) uniformly on the simplex
# -> theta0+theta1+theta2=1
theta = np.random.dirichlet((1, 1, 1), mini_batch_size)
```

We then generate data according to the random values of  $\theta$  above mentioned. We sample values from the corresponding 3-ary distribution and transform the result  $x \in \{0, 1, 2\}$  as an *one hot* vector.

```
# Generating data from distribution given by the 3-ary distribution
rand = np.random.uniform(0., 1., (mini_batch_size, 1))
cumul = np.cumsum(theta, axis=1)
data = 3 - np.sum(1.*(np.concatenate((rand, rand, rand), axis=1) < cumul),
                 axis=1)
```

```
# Change data as one_hot vector
one_hot_data = np.zeros((mini_batch_size,3))
one_hot_data[np.arange(mini_batch_size), data.astype(int)] = 1
```

Finally, we train the model to predict  $\theta$  from a single input.

```
_, train_mse = sess.run((train, mse),
                        feed_dict={x: one_hot_data, y_true: theta})
```

And test and display the results.

```
print('w: ', sess.run(w))
print('b: ', sess.run(b))
# estimated theta when 0 is observed
print('t0:', sess.run(y_pred, feed_dict={x:[[1,0,0]]}))
# estimated theta when 1 is observed
print('t1:', sess.run(y_pred, feed_dict={x:[[0,1,0]]}))
# estimated theta when 2 is observed
print('t2:', sess.run(y_pred, feed_dict={x:[[0,0,1]]}))
```

## 1.2 Outputs

```
=====
w: [[ 0.75430816 -0.26901755 -0.39179456]
     [-0.21355353  0.1496648  -0.6626014 ]
     [ 0.14008978 -0.18631633  0.38177642]]
b: [-0.15940666  0.17108977  0.2883165 ]
t0: [[0.50061333  0.25038624  0.24900039]]
t1: [[0.25001177  0.50030744  0.24968079]]
t2: [[0.25021031  0.25123584  0.49855384]]
test mse=0.0414
```

Figure I.1 – Outputs

The results (values of  $\theta$  and MSE) are similar to the theoretical values found in the next section (see subsection 2.4).

## 2 Derivations of the optimal estimator and the optimal average MSE

The goal is to determine :

$$\hat{\theta}_{MSE} = \operatorname{argmin}_{\hat{\theta}} \left[ 2 \int_0^1 \int_0^{1-\theta_0} \mathbb{E}_{X \sim p_{\theta}(x)} \left[ \|\hat{\theta}(X) - \theta\|^2 \right] d\theta_1 d\theta_0 \right]$$

We have :

$$\begin{aligned} \mathbb{E}_{X \sim p_{\theta}(x)} \left( \|\hat{\theta}(X) - \theta\|^2 \right) &= \mathbb{E}_{X \sim p_{\theta}(x)} \left( \sum_{i=0}^2 \left( \hat{\theta}_i(X) - \theta_i \right)^2 \right) \\ &= \sum_i \left( \mathbb{E} \left( [\hat{\theta}_i(X)]^2 \right) - 2\theta_i \mathbb{E} \left( \hat{\theta}_i(X) \right) + \mathbb{E} \left( \theta_i^2 \right) \right) \\ &= \underbrace{\sum_{i=0}^2 \sum_{j=0}^2 \theta_j [\hat{\theta}_i(j)]^2}_{=(1)} - 2 \underbrace{\sum_{i=0}^2 \sum_{j=0}^2 \theta_i \theta_j \hat{\theta}_i(j)}_{=(2)} + \underbrace{\sum_{i=0}^2 \theta_i^2}_{=(3)} \end{aligned}$$

Now, we can integrate each part of the MSE over  $\theta_0$  and  $\theta_1$ . Because the probability distribution of  $\theta = \{\theta_0, \theta_1, \theta_2\}$  is symmetric, we can determine  $\hat{\theta}(0) = [\hat{\theta}_0(0), \hat{\theta}_1(0), \hat{\theta}_2(0)]$  and infer the values of  $\hat{\theta}(1)$  and  $\hat{\theta}(2)$ .

### 2.1 Integration of the 1st part

Let's calculate :

$$\int_0^1 \int_0^{1-\theta_0} (1) d\theta_1 d\theta_0$$

with :

$$\begin{aligned} (1) &= \theta_0 \left( [\hat{\theta}_0(0)]^2 + [\hat{\theta}_1(0)]^2 + [\hat{\theta}_2(0)]^2 \right) + \theta_1 \left( [\hat{\theta}_0(1)]^2 + [\hat{\theta}_1(1)]^2 + [\hat{\theta}_2(1)]^2 \right) + \\ &\quad \theta_2 \left( [\hat{\theta}_0(2)]^2 + [\hat{\theta}_1(2)]^2 + [\hat{\theta}_2(2)]^2 \right) \end{aligned}$$

We find that :

$$\int_0^1 \int_0^{1-\theta_0} (1) d\theta_1 d\theta_0 = \sum_{j=0}^2 \frac{\|\hat{\theta}(j)\|^2}{6}$$

### 2.2 Integration of the 2nd part

Let's now calculate :

$$\int_0^1 \int_0^{1-\theta_0} (2) d\theta_1 d\theta_0$$

with :

$$\begin{aligned} (2) &= -2[\theta_0^2 \hat{\theta}_0(0) + \theta_0 \theta_1 \hat{\theta}_1(0) + \theta_0 \theta_2 \hat{\theta}_2(0) + \theta_0 \theta_1 \hat{\theta}_0(1) + \theta_1^2 \hat{\theta}_1(1) + \theta_1 \theta_2 \hat{\theta}_2(1) + \theta_0 \theta_2 \hat{\theta}_0(2) + \theta_1 \theta_2 \hat{\theta}_1(2) \\ &\quad + \theta_2^2 \hat{\theta}_2(2)] \end{aligned}$$

By keeping only the terms refering to  $\hat{\theta}(0)$ , we have:

$$-2 \int_0^1 \int_0^{1-\theta_0} \theta_0^2 \hat{\theta}_0(0) + \theta_0 \theta_1 \hat{\theta}_1(0) + \theta_0(1 - \theta_0 - \theta_1) \hat{\theta}_2(0) d\theta_1 d\theta_0 = -\frac{1}{6} \hat{\theta}_0(0) - \frac{1}{12} \hat{\theta}_1(0) - \frac{1}{12} \hat{\theta}_2(0)$$

Because the probability distribution of  $\theta$  is symetric, we can infer the results of the integral for the terms  $\hat{\theta}(1)$  and  $\hat{\theta}(2)$  and thus we find that:

$$\begin{aligned} \int_0^1 \int_0^{1-\theta_0} (2) d\theta_1 d\theta_0 &= -\frac{1}{6} \hat{\theta}_0(0) - \frac{1}{12} \hat{\theta}_1(0) - \frac{1}{12} \hat{\theta}_2(0) \\ &\quad - \frac{1}{12} \hat{\theta}_0(1) - \frac{1}{6} \hat{\theta}_1(1) - \frac{1}{12} \hat{\theta}_2(1) \\ &\quad - \frac{1}{12} \hat{\theta}_0(2) - \frac{1}{12} \hat{\theta}_1(2) - \frac{1}{6} \hat{\theta}_2(2) \end{aligned}$$

### 2.3 Integration of the 3rd part

Let's finally calculate :

$$\int_0^1 \int_0^{1-\theta_0} (3) d\theta_1 d\theta_0$$

with :

$$\begin{aligned} (3) &= \theta_0^2 + \theta_1^2 + \theta_2^2 \\ &= \theta_0^2 + \theta_1^2 + (1 - \theta_0 - \theta_1)^2 \\ &= 2\theta_0^2 + 2\theta_1^2 - 2\theta_0 - 2\theta_1 + 2\theta_0\theta_1 + 1 \end{aligned}$$

We find :

$$\begin{aligned} \int_0^1 \int_0^{1-\theta_0} 2\theta_0^2 + 2\theta_1^2 - 2\theta_0 - 2\theta_1 + 2\theta_0\theta_1 + 1 d\theta_1 d\theta_0 &= \int_0^1 \left( -\frac{5}{3} \theta_0^3 + 3\theta_0^2 - 2\theta_0 + \frac{2}{3} \right) d\theta_0 \\ &= \frac{1}{4} \end{aligned}$$

### 2.4 Minimization of average MSE over all $\theta$

Thus, the estimator based one sample that minimizes the average MSE over all  $\theta$  is given by:

$$\begin{aligned} \hat{\theta}_{MSE} &= \operatorname{argmin}_{\hat{\theta}} \left[ 2 \int_0^1 \int_0^{1-\theta_0} \mathbb{E}_{X \sim p_{\theta}(x)} \left[ \|\hat{\theta}(X) - \theta\|^2 \right] d\theta_1 d\theta_0 \right] \\ &= \operatorname{argmin}_{\hat{\theta}} \left[ 2 \int_0^1 \int_0^{1-\theta_0} (1) + (2) + (3) d\theta_1 d\theta_0 \right] \\ &= \operatorname{argmin}_{\hat{\theta}} \left[ \begin{aligned} &\frac{1}{8} + \frac{1}{6} \left( \hat{\theta}_0(0) - \frac{1}{2} \right)^2 + \frac{1}{6} \left( \hat{\theta}_1(0) - \frac{1}{4} \right)^2 + \frac{1}{6} \left( \hat{\theta}_2(0) - \frac{1}{4} \right)^2 \\ &\frac{1}{6} \left( \hat{\theta}_0(1) - \frac{1}{4} \right)^2 + \frac{1}{6} \left( \hat{\theta}_1(1) - \frac{1}{2} \right)^2 + \frac{1}{6} \left( \hat{\theta}_2(1) - \frac{1}{4} \right)^2 \\ &\frac{1}{6} \left( \hat{\theta}_0(2) - \frac{1}{4} \right)^2 + \frac{1}{6} \left( \hat{\theta}_1(2) - \frac{1}{4} \right)^2 + \frac{1}{6} \left( \hat{\theta}_2(2) - \frac{1}{2} \right)^2 \end{aligned} \right] \end{aligned}$$

Finally,

$$\hat{\theta}_{MSE}(x) = \begin{cases} (1/2, 1/4, 1/4) & \text{if } x = 0 \\ (1/4, 1/2, 1/4) & \text{if } x = 1 \\ (1/4, 1/4, 1/2) & \text{if } x = 2 \end{cases}$$

For this estimator, the MSE averaged over all  $\theta$  is equal to  $\frac{1}{8} = 3 \times \frac{1}{24}$ . Thus, the error contribution for each  $\hat{\theta}_i$ ,  $i \in \{0, 1, 2\}$  is  $\frac{1}{24} \approx 0.04$ . We find the same result than the practical result (section [1.2](#)).

# Part II

## Task 2 : Rain or shine

### 1 Explanation and outputs of *project1\_task2\_generate.py*

#### 1.1 Explanation

Function to approximate the steady-state by multiplying the transition matrix of the markov chain. The stop condition is  $|P^{i+1} - P^i| < \epsilon$  or  $i > 10000$

```
# Approximate the steady-state by multiplying the transition
# matrix of the markov chain until convergence.
def approx_steady(t_matrix, eps=0.0001):
    def dist_max(x,y):
        return max(np.abs(x-y))
    i = 0
    t0 = t_matrix
    t1 = np.matmul(t0, t_matrix)
    while (i<10000 and dist_max(t0[0],t1[0])>eps):
        t0=t1
        t1 = np.matmul(t0, t_matrix)
    return t0[0]
```

The function *generate\_next\_state* generates a next state  $X_{t+1}$  given previous state  $X_t$  and transition matrix  $P$ .

```
# Return a state given a list of proba :
#t_prob = [p(0), p(1), p(2), p(3)]
def generate_state(t_prob):
    p = np.random.rand(1)
    cumsum = np.cumsum(t_prob)
    for i, val in enumerate(cumsum):
        if (p<val):
            return i

# Generate next state given previous state
# and transition matrix
def generate_next_state(prev_state, tr_mat):
    return generate_state(tr_mat[prev_state])
```



We then can generate easily  $n = 100000$  data given the transition matrix describes by  $psh, pra, pcl, psn$ .  $X_0$  is generated thanks to the approximation of the steady-state probabilities.

```
# prob. conditional of sh (shine), ra (rain), cl (cloudy),
# sn (snow) given for each shine, rain, cloudy, snow
psh = [0.6, 0.1, 0.1, 0.1]
pra = [0.1, 0.5, 0.3, 0.1]
pcl = [0.2, 0.3, 0.4, 0.3]
psn = [0.1, 0.1, 0.2, 0.5]

# create transition matrix of the markov chain
tr_mat = np.transpose([psh, pra, pcl, psn])

# steady-state prob. of respectively shine, rain, cloudy, snow
pr = approx_steady(tr_mat)
print("Approx steady-state prob. : {}".format(pr))

n = 100000
x = np.zeros(n, dtype=np.int)

for k in range(n):
    if k == 0:
        x[k] = generate_state(pr)
    else:
        x[k] = generate_next_state(x[k-1], tr_mat)
```

## 1.2 Outputs

```
Approx steady-state prob. : [0.20019531 0.27027365 0.31106228 0.21846876]
> 19.982% shine days in data
> 27.11% rain days in data
> 31.178% cloudy days in data
> 21.73% snow days in data
```

Figure II.1 – Stats of data

## 2 Explanation and outputs of *project1\_task2.py*

### 2.1 Explanations

We assume in this example that the weather at day  $t$  is fixed. It can take only one value given by shine (0), rain (1), cloudy (2), snow (3). Thus, we create a placeholder for the input (day  $t$ ) and one for the label (day  $t+1$ ) which are integers  $x \in \{0, 1, 2, 3\}$  and translate these values in "one hot" vector to help our neural network training.

```
# Create placeholder of input : the weather at day t.
# shine (0), rain (1), cloudy (2), snow (3)
x = tf.placeholder(tf.int32, shape=[None, 1])
x_hot = tf.one_hot(x, 4)
```

```
# Create placeholder of input : the weather at day t+1.
# shine (0), rain (1), cloudy (2), snow (3)
y_true = tf.placeholder(tf.int32, shape=[None, 1])
y_true_hot = tf.one_hot(y_true, 4)
```

Then, we create our neural network architecture. It consists of a simple multinomial logistic regression with four inputs and four outputs. Thus, the computed function is :

$$p(y = i | x; w, b) = \text{softmax}(w^T x + b)_i, \quad \forall i \in \{0, 1, 2, 3\}$$

with  $x \in \mathbb{R}^4$ ,  $w \in \mathbb{R}^4 \times \mathbb{R}^4$  and  $b \in \mathbb{R}^4$

```
w = tf.Variable(tf.truncated_normal([4, 4], stddev=0.5))
b = tf.Variable(tf.constant(0.1, shape=[4]))
z = tf.matmul(x_hot, w) + b
y_pred = tf.nn.softmax(z)

loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=y_true_hot,
                                              logits=z))
train = tf.train.AdamOptimizer(0.001).minimize(loss)
```

We train the model as for the rain and shine problem.

```
sess = tf.InteractiveSession()
sess.run(tf.global_variables_initializer())
for epoch in range(30):
    rindex = np.random.permutation(n)
    s = s[rindex,:]
    p = p[rindex,:]
    for i in range(n // m):
        sess.run(    train,
                    feed_dict={
                        x: s[i*m:i*m+m,:],
                        y_true: p[i*m:i*m+m,:]
                    })
```

Finally, we test the model and display results. We display the probability of weather at day  $t + 1$  given weather at day  $t$  for all possible weather at day  $t$ .

```
s_test = np.zeros((4, 1))
s_test[0] = 0
s_test[1] = 1
s_test[2] = 2
s_test[3] = 3

# Evaluate prediction for each weather at day t
pred = y_pred.eval(feed_dict={x: s_test})

# Display results
print()
print("——— Resultats pred prob. cond ———")
```

```

print("Proba[ X(t+1) | X(t) = 'shine' ] = {}".format(pred[0]))
print("Proba[ X(t+1) | X(t) = 'rain' ] = {}".format(pred[1]))
print("Proba[ X(t+1) | X(t) = 'cloudy' ] = {}".format(pred[2]))
print("Proba[ X(t+1) | X(t) = 'snow' ] = {}".format(pred[3]))

```

## 2.2 Outputs

```

----- Resultats pred prob. cond -----
Proba[ X(t+1) | X(t) = 'shine' ] = [[0.59664243 0.10369433 0.20219867 0.09746451]]
Proba[ X(t+1) | X(t) = 'rain' ] = [[0.09989394 0.49063125 0.30890697 0.1005678 ]]
Proba[ X(t+1) | X(t) = 'cloudy' ] = [[0.09801193 0.29907972 0.40050235 0.20240599]]
Proba[ X(t+1) | X(t) = 'snow' ] = [[0.10040201 0.10248841 0.30068332 0.49642625]]

```

Figure II.2 – Prediction prob. of weather at day t+1 given weather at day t

We can see that the neural network have approximately learnt the conditional probabilities given in the table of proba. There is a small error probably due to the small number of initial data.